GPT Pedagogy

Matthew Pisano[1], Tiburon Benavides[1], Lu Zhou[1], Yanshen Lin[1], and Yiyang Cai[1]

[1]Rensselaer Polytechnic Institute

April 16, 2023

Abstract

GPT-3 is a Large Language Model which, through overuse, is often used to deprive students of the opportunity to learn effectively. Our intended use case of GPT enables the transparent use of the technology between instructor and student. We aim to create a more active and participatory learning environment through the usage of the model in active learning. Our long-term goal for higher education is along the lines of the fourth UN SDG:

> Ensure inclusive and equitable quality education and promote lifelong learning opportunities for all.

We view GPT as a way to create an adaptive learning experience which promotes the educational endeavor rather than detracts from it.

Our plan is to develop a GPT-3 based learning assistant for the *Introduction to Biology* course here at RPI. Through a partnership with the professor, we will gather sufficient amounts of course material to fine-tune a pre-trained GPT-3 model. This will create a knowledgeable and focused learning assistant, *Mathesis* (from the ancient greek work for learning). By creating this personalized AI tutor we can make it easier and more transparent for students to learn the material and reduce their stress.

One of our goals is for the model to maintain its conversational abilities while embedding additional knowledge about faculty defined key learning objectives. *Mathesis* will generate a series of topic-relevant questions, evaluate the answers of those questions, and give useful feedback or counter-examples to the student. The model benefits from human-in-the-loop reinforcement learning by storing previous chats from students and faculty alike.

We also aim for this model to be generalizable to other classes and disciplines in the future. The model will work especially well for courses where it is difficult to give personalized feedback to each learner in each class meeting time, such as in classes that have a high student to faculty ratio. It will also work well for courses with students interested in AI who cannot adequately engage with that interest through the course material.

1 Introduction

The idea for this project was conceived before the beginning of this project in the context of our Informatics class. It stemmed from observations about the current trend that large language models are heading in and how many educational institutions are not fully prepared to handle the fallout of this trend. Our primary motivation for this project is to remedy this issue. By providing students with an opportunity to use large language models, like GPT, as a learning tool, we allow students to gain valuable insight into the practical benefits and shortcomings of these models.

1.1 Our Targeted Approach

When looking at our project through a more focused lens, we believe that it will benefit the students that take the into courses that we hope *Mathesis* will be integrated into. For many introductory classes, their main focus is to imbue students with a large amount of foundational knowledge before they explore specific categories further in subsequent classes. We believe that a GPT model will have a significant impact in helping with this task. By fine-tuning a model on a large amount of course material, it will be able to provide relevant questions from students with greater detail and with a stronger relation to the other topics in the course. This comes in addition to the large amount of general knowledge that the model initially has. While it may not be able to go into detail about very niche topics, fine-tuning allows it to provide adequate detail on the material that students need to learn.

1.2 Broad Impact

From a broader perspective, customized learning assistants, similar to ours, have the potential to make a significant impact on the educational landscape as a whole. One of the biggest issues that our project address is the asymmetric distribution of time between students and teachers. Many courses at RPI, especially one or two thousand level courses, have a very high student to faculty ratio. Due to this, each student is only afforded a very limited amount of one-on-one time with the professor or their TAs. Automated learning assistants like *Mathesis* can help to afford students more time with a knowledgeable entity in their field. This effect would not only benefit students at RPI, but also nearly every other student in every other institution. While models such as ours cannot serve as a complete replacement for faculty, they can instead serve as an aide, answering students' questions and providing helpful feedback. The more competent this model, the less one-on-one time students will need to spend with faculty. This will hopefully result in a more complete understanding of the course material among the students.

1.3 Aristotle For All

This has served as a large motivator for us while working on this project. The opportunity to both free up faculty to have more personalized time with each student and to allow each student to have a personal tutor by their side will help to solve many of the struggles that students have with learning in general. These ideas are reflected in both past and contemporary sources. In a very similar vein to our project, Khan Academy's *Khanmigo*[1] model aims to provide a very similar service to students through the internet. Our hope is that *Mathesis* will soon serve

a similar role, but customized for RPI's students and curriculum. These ideas are reflected in the well researched, and somewhat prophetic, video *Digital Aristotle: Thoughts on the Future of Education* by CGP Grey. It outlines many of the same points that we have mentioned here in addition to detailing how technologies such as *Mathesis* or *Khanmigo* can have on the future of education.

2   Project Use Case

2.1   Specifications

For the development of the use case for this project, it is important to detail several specifications that describe GPT Pedagogy in a more rigorous manner. One of these specifications is the list of the requirements of the project. To fulfil our objectives, the project has the following *functional requirements*:

- Provide students with a knowledgeable teacher that can answer any course-relevant question

- Be able to automatically generate a series of evaluation questions for students and reference answers based on the core topics of the course

- Receive students' answers to evaluation questions and evaluate their correctness based off of the reference answers

- Provide useful feedback to students if their answers did not sufficiently match the reference answers

- Allow instructors to review and regenerate all questions generated by the model

- Allow instructors to evaluate a summary of student performance based off of the automatically generated questions

Along with these functional requirements, the project requires several *non-functional requirements*. Our project must also have:

- A GPT model that can respond in a conversational and contextually sensitive manner

- That model be fine-tuned on relevant course material to narrow its focus of expertise

- A front end environment that allows students and administrators to interact with the model

- A login and authentication system to prevent unwanted access and to keep track of evaluations

Another important specification to consider is that of the entropy/uncertainty of the project. Overall, we have worked to minimize the total uncertainty of the project. It is important to note, however, that the GPT model will always introduce some amount of extra uncertainty.

One of the areas where we worked to minimize the uncertainty in was the user interface. This is split into two parts, the student and administrator views. In both of these cases, our design was oriented towards a simple, straightforward interface. We worked to accomplish this through the use of a minimalist design. Users can choose to choose only a few tabs: the main chat and the lesson evaluations. The administrators have a similar view, with th addition of the ability to edit the

questions that are displayed to students. In our design, buttons, animations, and images are kept to a minimum. This will hopefully allow users to focus more on working on evaluations or interacting with *Mathesis* to better learn specific topics.

Another way we worked to minimize uncertainty is through the fine-tuning of the model on our own, customized, training data. By feeding *Mathesis* significant amounts of course material, we have narrowed its focus down to the relevant topics of the course. While the model does not forget its previously learned knowledge, its ability to correctly interpret and recall information related to the course has increased. This lowers the uncertainty of the project by encouraging the trained model to generate responses that are targeted towards the course material. This will increase its helpfulness to students as unrelated responses may provide misinformation or serve to demotivate students from interacting with the system.

The use case that we developed addresses what we expect to be the basic low of the system, along with any reasonable alternate or exception flows. The goal of this system is the same as the project in general: to use an interface and pre-trained model to provide students with a flexible and helpful learning assistant for the *Introduction to Biology* course at RPI.

We have included the particular flows that we did in order to properly scope our use case. Our intention is to cover all relevant ways in which students and administrators, while not including flows that may bloat our design of this early use case. Examples of such flows would be deliberate adversarial attacks to the system. While these attacks are a near certainty in a deployed system, the time constraints on this project has not allowed us to account for them in our working prototype. Due to this reason, we have set this and similar edge cases to be outside our scope.

2.2   Implementation Considerations

As mentioned above, one of our primary design goals of this project was to minimize uncertainty. In addition to the previously mentioned methods, we also minimized uncertainty through our usage of semiotics, cognitive principles, and the project architecture.

One of the ways in which we minimize uncertainty is through our placement of our signs. Many of the elements in the user interface lack symbols or icons. This is by design. As this project is currently in a prototype form, it would be unwise to begin using signs in places where they do not need to be. In place of these signs we instead have text descriptions. While this may be less visually appealing, it offers a lower uncertainty as many elements have their descriptions written on them. An example of a sign that we do use can be found in the 'send' icon for the chat functionality. This sign is an icon, specifically as it does not physically resemble anything being sent. We chose this icon to be that of a paper airplane, as is standard across most massaging applications. Through the use of this, and other icons, we hope to make all functionality clear while maintaining familiarity.

We also use both semiotics and cognitive principles in the chat and feedback functionalities of the project. Our design of the chat and feedback interfaces mirrors that of more mundane chat-focused applications. By modeling our interface in a similar manner to SMS communications, we create a familiar environment to both students and teachers. This was likely the rationale of ChatGPT's interface design, which mirrors these communications as well. The lesson quiz submitting also mirrors familiar environments such as LMS. By creating a familiar environment, we hope that users will be able to immediately recognize the form and function of

the interface, allowing for quicker acclimation.

3   Architectural Models

3.1   Conceptual Model

...

3.2   Logical Model

...

3.3   Implementation

...

4   Prototype Implementation

Over the course of the project, we made several important decisions on how to best implement our prototype design.

4.1   Model Choice

One of the biggest references that we used in this decision-making process was the OpenAi Python API documentation[2]. This documentation details each of the models that we can use, their capabilities, and their drawbacks. On example of this comes in the type of model that we use. As our primary model, we use their *text-davinci-003* model. We have found that this model is a good balance between cost, quality, and flexibility. DaVinci is the most advances of the *text-** models. It offers similar quality to ChatGPT and has the added bonus of being available for file-tuning. Using this feature helps to lower the uncertainty of our information system, as mentioned above. An option that we did not decide to use as our primary model is the *gpt-3.5-turbo*. This model is the model behind ChatGPT itself; it is more expensive and has the drawback of not being pre-trainable. Due to these factors, we decided to keep using *text-davinci-003* after the announcement of *gpt-3.5-turbo* that came mid-way through the project. This decision helps us to fulfil the first four of our functional requirements.

4.2   Alternate Flows

Another important decision that we made was our decision to exclude particular alternative or error flows from our design. From the beginning of this project, we have been aware that large language models like the GPT series can both be incredibly helpful and somewhat volatile. Notably, we considered the possibility that users could perform adversarial attacks on the model. While these attacks would only affect that particular student, it is important to make sure our model provides all students with an appropriate learning environment, regardless of their actions. This decision was influenced primarily by the *DAN*[3] (Do Anything Now) exploit and its many iterations. This exploit allows users to convince the ChatGPT model to ignore its alignment and safety training. Attacks like these would likely

also be able to bypass the fine-tuning of *Mathesis*. We have partially confirmed this assumption by testing similar exploits on GPT models. We performed this evaluation by taking note of several topics that GPT disliked talking about and convincing the model to mention them, in our case, past political figures, by speaking about adjacent topics then re-asking questions directing at these disallowed topics. Our evaluation suggests that students could attempt something similar if they wished to do so. At present, we have decided to not work on measures specifically preventing this type of attack. The reason for this is the relatively quick patches OpenAI sends out when exploits are discovered combined with the difficulty of the problem. Our hope is that this will not impact the functional requirements of the project, especially as OpenAI continues to work against the issue.

5   Conclusion

5.1   Future Work

In our future work in this project, we expect to both finalize the attributes of the project that we have included in its current scope and to expand the projects scope to cover more edge cases.

...

Ways in which we would expand the scope of the project would to be to better handle reasonable edge cases. These could come in the form of students needing help with navigating the new system, handling adversarial attacks from students, or further automating the process of training models on new datasets. These possible additions would serve to help students better adapt to using the *Mathesis* assistant and its surrounding learning platform, deter students from using the model for unintended purposes, and creating encouraging environment for other professors who would like to include their courses into the learning system in future semesters respectively.

5.2   Closing Thoughts

...

References

[1] K. Academy. Khanmigo education ai guide. `https://www.khanacademy.org/khan-labs`. Accessed: 2023-04-16.

[2] OpenAI. Introduction - openai documentation. `https://platform.openai.com/docs/introduction`. Accessed: 2023-04-16.

[3] u/walkerspider. Dan is my new friend. `https://www.reddit.com/r/ChatGPT/comments/zlcyr9/dan_is_my_new_friend/`. Accessed: 2023-04-16.